

Buffer overrun attacks

Myrto Arapinis
School of Informatics
University of Edinburgh

Basic buffer overflow (demo recap)

```
#include <stdio.h>
#include <stdlib.h>

int read_stdin(void){
    char buf[128];
    int i;
    gets(buf);
    i = atoi(buf);
    return i;
}

int main(int ac, char **av){
    int x = read_stdin();
    printf("x=%d\n", x);
}
```

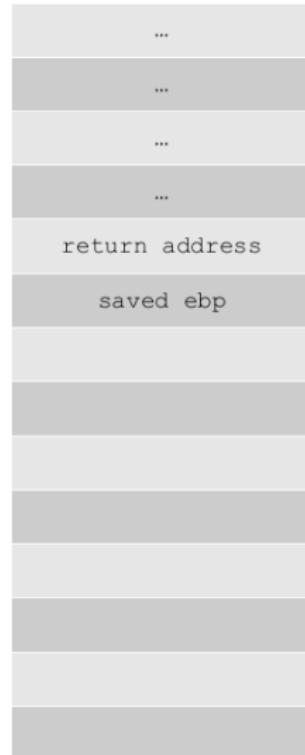


Basic buffer overflow (demo recap)

```
#include <stdio.h>
#include <stdlib.h>

int read_stdin(void){
    char buf[128];
    int i;
    gets(buf);
    i = atoi(buf);
    return i;
}

int main(int ac, char **av){
    int x = read_stdin();
    printf("x=%d\n", x);
}
```

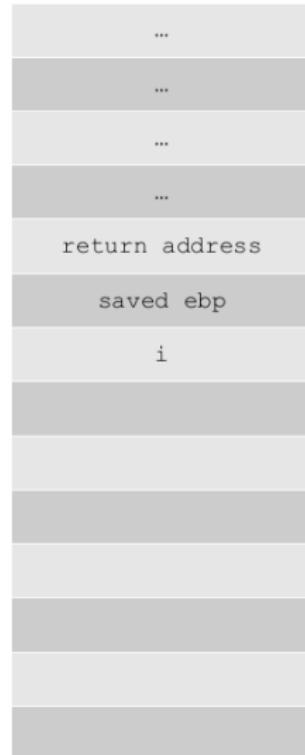


Basic buffer overflow (demo recap)

```
#include <stdio.h>
#include <stdlib.h>

int read_stdin(void){
    char buf[128];
    int i;
    gets(buf);
    i = atoi(buf);
    return i;
}

int main(int ac, char **av){
    int x = read_stdin();
    printf("x=%d\n", x);
}
```



Basic buffer overflow (demo recap)

```
#include <stdio.h>
#include <stdlib.h>

int read_stdin(void){
    char buf[128];
    int i;
    gets(buf);
    i = atoi(buf);
    return i;
}

int main(int ac, char **av){
    int x = read_stdin();
    printf("x=%d\n", x);
}
```



Basic buffer overflow (demo recap)

```
#include <stdio.h>
#include <stdlib.h>

int read_stdin(void){
    char buf[128];
    int i;
    gets(buf);
    i = atoi(buf);
    return i;
}

int main(int ac, char **av){
    int x = read_stdin();
    printf("x=%d\n", x);
}
```

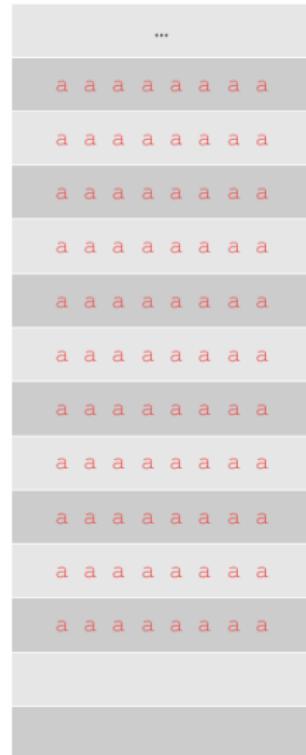


Basic buffer overflow (demo recap)

```
#include <stdio.h>
#include <stdlib.h>

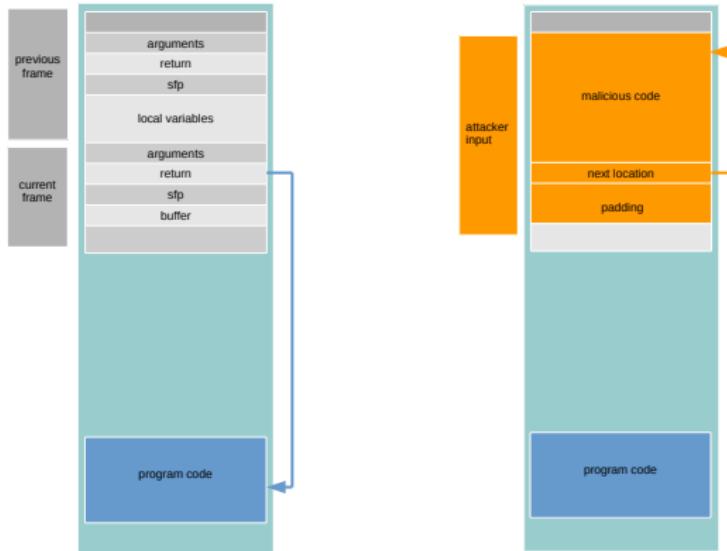
int read_stdin(void){
    char buf[128];
    int i;
    gets(buf);
    i = atoi(buf);
    return i;
}

int main(int ac, char **av){
    int x = read_stdin();
    printf("x=%d\n", x);
}
```



Exploiting buffer overruns

Control hijacking



A buffer overflow can change the flow of execution of the program:

- ▶ load malicious code into memory
- ▶ make %eip point to it

Shellcode injection

Goal: “spawn a shell” - will give the attacker general access to the system

```
#include stdio.h
void main() {
char *name[2];
name[0] = "/bin/sh";
name[1] = NULL;
execve(name[0], name, NULL);
}
```

C code

```
"\x31\xc0"
"\x50"
"\x68" "//sh"
"\x68" "/bin"
"\x89\xe3"
"\x50"
...
```

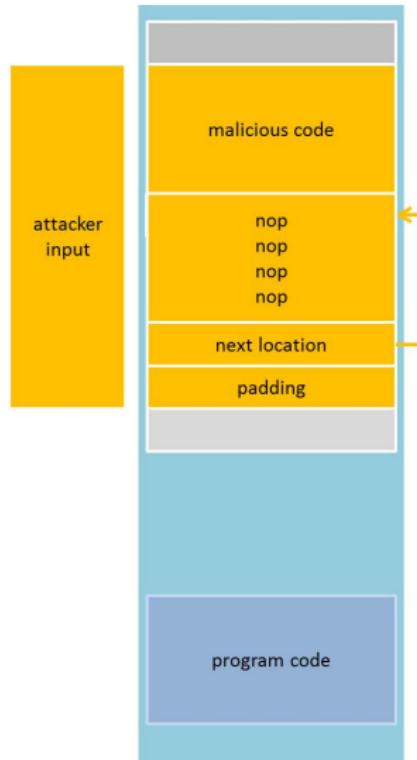
Machine code
(part of attacker's input)

- ▶ must inject the machine code instructions (code ready to run)
- ▶ the code cannot contain any zero bytes (`printf`, `gets`, `strcpy` will stop copying)
- ▶ can't use the loader (we're injecting)

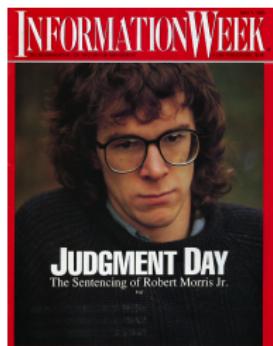
The return address

Challenge: find the address of the injected malicious code?

- ▶ If code accessible: we know how far is the overflowed variable from the saved %ebp
- ▶ If code not accessible: try different possibilities!
In a 32 bits memory space, there are 2^{32} possibilities
- ▶ NOP sled
 - ▶ guess approximate stack state when the function is called
 - ▶ insert many NOPs before Shell Code

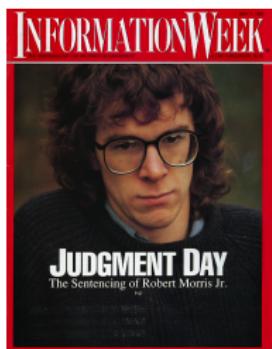


First exploits of BOs



1988 - Robert Morris Jr.

First exploits of BOs



1988 - Robert Morris Jr.

Title : Smashing The Stack For Fun And Profit

Author : Aleph1

.00 Phrack 49 0o.

Volume Seven, Issue Forty-Nine

File 14 of 16

BugTraq, r0t, and Underground.org
bring you

XXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXX
Smashing The Stack For Fun And Profit
XXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXX

by Aleph One
aleph1@underground.org

'smash the stack' [C programming] n. On many C implementations it is possible to corrupt the execution stack by writing past the end of an array declared auto in a routine. Code that does this is said to smash the stack, and can cause return from the routine to jump to a random address. This can produce some of the most interesting and dangerous bugs known to mankind. Variants include trash the stack, scribble the stack, mangle the stack; the term nung the stack is not used, as this is never done intentionally. See spm; see also alias bug, fandango on core, memory leak, precedence lossage, overrun screw.

Introduction

Over the last few months there has been a large increase of buffer overflow vulnerabilities being both discovered and exploited. Examples of these are syslog, splitvt, library, at, etc. This paper are, and how their exploits work.

Basic knowledge of assembly memory concepts, and experience. We also assume we are working system is Linux.

Some basic definitions before block of computer memory that type. C programmers normally commonly, character arrays. A declared either static or dynamic time on the data segment. Dynamic the stack. To sacrifice is to give. We will concern ourselves only known as stack-based buffer overflows.



1996 - Elias Levy (Aleph One)

Buffer overrun opportunities

Unsafe libc functions

```
strcpy (char *dest, const char *src)
```

```
strcat (char *dest, const char *src)
```

```
gets (char *s)
```

```
scanf (const char *format, ...)
```

```
...
```

Do not check bounds of buffers they manipulate!!

Use safe functions

```
int read_stdin(void){  
    char buf[128];  
    int i;  
    fgets(buf, sizeof(buf), stdin);  
    i = atoi(buf);  
    return i;  
}
```

But then...

... your program is as secure as its programmer is cautious. It is now up to the programmer to include all the necessary checks in his program :-/ and this is a tricky one...

Arithmetic overflows

- ▶ Limitation related to the representation of integers in memory
- ▶ In 32 bits architectures, signed integers are expressed in **two's compliment notation**
 - $0x00000000$ - $0x7fffffff$: positive numbers $0 - (2^{31} - 1)$
 - $0x80000000$ - $0xffffffff$: negative numbers $(-2^{31} + 1) - (-1)$
- ▶ In 32 bits architectures, unsigned integers are only positive numbers $0x00000000$ - $0xffffffff$.
Once the highest unsigned integer is reached, the next sequential integer wraps around zero.

```
# include <stdio.h>
int main(void){
    unsigned int num = 0xffffffff;
    printf("num + 1 = 0%x\n", num + 1);
    return 0;
}
```

The output of this program is: **num + 1 = 0x0**

Arithmetic overflow exploit (1)

- ▶ Stack-based buffer overflow due to arithmetic overflow:

```
int catvars(char *buf1, char *buf2,
            unsigned int len1, unsigned int len2){
    char mybuf[256];
    if((len1 + len2) > 256){
        return -1;
    }
    memcpy(mybuf, buf1, len1);
    memcpy(mybuf + len1, buf2, len2);
    do_something(mybuf);
    return 0;
}
```

Arithmetic overflow exploit (1)

- ▶ Stack-based buffer overflow due to arithmetic overflow:

```
int catvars(char *buf1, char *buf2,
            unsigned int len1, unsigned int len2){
    char mybuf[256];
    if((len1 + len2) > 256){
        return -1;
    }
    memcpy(mybuf, buf1, len1);
    memcpy(mybuf + len1, buf2, len2);
    do_something(mybuf);
    return 0;
}
```

Check can be bypassed by using suitable values for len1 and len2, e.g. len1 = 0x00000103, len2 = 0xfffffffffc, len1+len2 = 0x000000ff (decimal 255)

Arithmetic overflow exploit (2)

- ▶ Stack-based buffer overflow due to arithmetic overflow:

```
int catvars(char *buf, int len){  
    char mybuf[256];  
    if(len > 256){  
        return -1;  
    }  
    memcpy(mybuf, buf, len);  
    return 0;  
}
```

Arithmetic overflow exploit (2)

- ▶ Stack-based buffer overflow due to arithmetic overflow:

```
int catvars(char *buf, int len){  
    char mybuf[256];  
    if(len > 256){  
        return -1;  
    }  
    memcpy(mybuf, buf, len);  
    return 0;  
}
```

```
memcpy(void *s1, const void *s2, size_t n); // size_t  
is unsigned
```

Check can be bypassed by using suitable values for len,
e.g. len = -1 = 0xffffffff , will be interpreted as an
unsigned integer encoding the value $2^{32} - 1!$

Arithmetic overflow exploit (3)

- ▶ Heap-based buffer overflow due to arithmetic overflow:
 - Memory **dynamically** allocated will persist across multiple function calls.
 - This memory is allocated on the **heap** segment.
 - Heap-based buffer overflows are more complex, and require understanding garbage collection and heap implementation.

Arithmetic overflow exploit (3)

- ▶ Heap-based buffer overflow due to arithmetic overflow:
 - Memory **dynamically** allocated will persist across multiple function calls.
 - This memory is allocated on the **heap** segment.
 - Heap-based buffer overflows are more complex, and require understanding garbage collection and heap implementation.

```
int myfunction(int *array, int len){  
    int *myarray, i;  
    myarray = malloc(len * sizeof(int));  
    if(myarray == NULL){  
        return -1;  
    }  
    for(i = 0; i < len; i++){  
        myarray[i] = array[i];  
    }  
    return myarray;  
}
```

Arithmetic overflow exploit (3)

- ▶ Heap-based buffer overflow due to arithmetic overflow:
 - Memory **dynamically** allocated will persist across multiple function calls.
 - This memory is allocated on the **heap** segment.
 - Heap-based buffer overflows are more complex, and require understanding garbage collection and heap implementation.

```
int myfunction(int *array, int len){  
    int *myarray, i;  
    myarray = malloc(len * sizeof(int));  
    if(myarray == NULL){  
        return -1;  
    }  
    for(i = 0; i < len; i++){  
        myarray[i] = array[i];  
    }  
    return myarray;  
}
```

Can allocate a size 0 buffer for myarray by using suitable value for len: len = 1073741824 , sizeof(int) = 4,
len*sizeof(int) = 0

The Ariane 5 Disaster

[Blexim] Basic Integer Overflows <http://phrack.org/issues/60/10.html#article>

Attempt to store a value in an integer which is greater than the maximum value the integer can hold
→ the value will be truncated



Ariane 5 rocket launch explosion due to integer overflow

Format strings

[Ref] scut/team teso. Exploiting Format String Vulnerabilities

- ▶ A format function takes a variable number of arguments, from which one is the so called format string

Examples: `fprintf`, `printf`, ..., `syslog`, ...

```
printf("The amount is %d pounds\n", amnt);
```

Format strings

[Ref] scut/team teso. Exploiting Format String Vulnerabilities

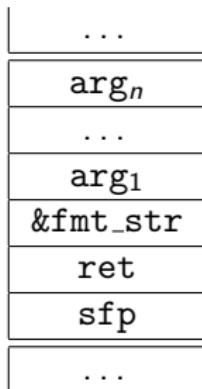
- ▶ A format function takes a variable number of arguments, from which one is the so called format string

Examples: `fprintf`, `printf`, ..., `syslog`, ...

`printf("The amount is %d pounds\n", amnt);`

- ▶ The behaviour of the format function is controlled by the format string. The function retrieves the parameters requested by the format string from the stack

Example: `printf(fmt_str, arg1, ..., argn);`



Exploiting format strings

- ▶ If an attacker is able to provide the format string to a format function, a format string vulnerability is present

```
int vulnerable_print(char *user) {  
    printf(user);  
}
```

```
int safe_print(char *user){  
    printf ("%s", user);  
}
```

Format strings exploits

- ▶ We can view the stack memory at any location
 - ▶ walk up stack until target pointer found
 - ▶ `printf ('%08x.%08x.%08x.%08x.%08x|%s|') ;`
 - ▶ A vulnerable program could leak information such as passwords, sessions, or crypto keys

- ▶ We can write to any memory location
 - ▶ `printf('hello %n', &temp)` – writes '6' into temp
 - ▶ `printf('hello%08x.%08x.%08x.%08x.%n')`

More buffer overflow opportunities

- ▶ Exception handlers
- ▶ Function pointers
- ▶ Double free
- ▶ ...

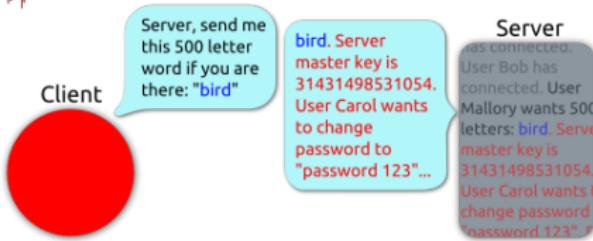


TLS Heartbleed

Heartbeat – Normal usage



Heartbeat – Malicious usage





TLS Heartbleed

Then, OpenSSL will uncomplainingly copy 65535 bytes from your request packet, even though you didn't send across that many bytes:

```
1  /* Allocate memory for the response, size is 1 byte
2   * message type, plus 2 bytes payload length, plus
3   * payload, plus padding
4   */
5   buffer = OPENSSL_malloc(1 + 2 + payload + padding);
6   bp = buffer;
7
8   /* Enter response type, length and copy payload */
9   *bp++ = TLS1_HB_RESPONSE;
10  s2n(payload, bp);
11  memcpy(bp, pl, payload);
12  bp += payload;
13  /* Random padding */
14  RAND_pseudo_bytes(bp, padding);
15
16  r = dtls1_write_bytes(s, TLS1_RT_HEARTBEAT, buffer, 3 + payload +
```

That means OpenSSL runs off the end of your data and scoops up whatever else is next to it in memory at the other end of the connection, for a potential data leakage of approximately 64KB each time you send a malformed heartbeat request.



TLS Heartbleed

How I used Heartbleed to steal a site's private crypto key | Ars Technica

https://tbb.../download/ Assessment ... Edinburgh NyMi Band E...ram | NyMi BufferOverflows ▾ SpoofingEmail ▾ ComputerSecurity ▾ SimSec ▾ How I used Heartbleed to steal a site's private crypto key | Ars Technica How Heartbleed Works: The Code Behind the

ars TECHNICA SEARCH BIZ & IT TECH SCIENCE POLICY CARS GAMING & CULTURE FORUMS SIGN IN ▾

BIZ & IT —

How I used Heartbleed to steal a site's private crypto key

Extracting keys from unpatched servers requires skill, but it's eminently doable.

RUBIN XU - 4/27/2014, 9:10 PM



/ Thinkstock

One of the most common attacks on memory safety

The CWE Top 25

Below is a brief listing of the weaknesses in the 2020 CWE Top 25, including the overall score of each.

Rank	ID	Name	Score
[1]	CWE-79	Improper Neutralization of Input During Web Page Generation ('Cross-site Scripting')	46.82
[2]	CWE-787	Out-of-bounds Write	46.17
[3]	CWE-20	Improper Input Validation	33.47
[4]	CWE-125	Out-of-bounds Read	26.50
[5]	CWE-119	Improper Restriction of Operations within the Bounds of a Memory Buffer	23.73
[6]	CWE-89	Improper Neutralization of Special Elements used in an SQL Command ('SQL Injection')	20.69
[7]	CWE-200	Exposure of Sensitive Information to an Unauthorized Actor	19.16
[8]	CWE-416	Use After Free	18.87
[9]	CWE-352	Cross-Site Request Forgery (CSRF)	17.29
[10]	CWE-78	Improper Neutralization of Special Elements used in an OS Command ('OS Command Injection')	16.44
[11]	CWE-190	Integer Overflow or Wraparound	15.81
[12]	CWE-22	Improper Limitation of a Pathname to a Restricted Directory ('Path Traversal')	13.67
[13]	CWE-476	NULL Pointer Dereference	8.35
[14]	CWE-287	Improper Authentication	8.17
[15]	CWE-434	Unrestricted Upload of File with Dangerous Type	7.38
[16]	CWE-732	Incorrect Permission Assignment for Critical Resource	6.95
[17]	CWE-94	Improper Control of Generation of Code ('Code Injection')	6.53
[18]	CWE-522	Insufficiently Protected Credentials	5.49
[19]	CWE-611	Improper Restriction of XML External Entity Reference	5.33
[20]	CWE-798	Use of Hard-coded Credentials	5.19
[21]	CWE-502	Deserialization of Untrusted Data	4.93
[22]	CWE-269	Improper Privilege Management	4.87
[23]	CWE-400	Uncontrolled Resource Consumption	4.14
[24]	CWE-306	Missing Authentication for Critical Function	3.85
[25]	CWE-862	Missing Authorization	3.77